



Homework 03

Name: McClain Thiel

Student ID: 3034003600

Note: Please print the output of each question in a new cell below your code

Numpy Introduction

1a) Create two numpy arrays (a and b). a should be all integers between 25-35 (inclusive), and b should be ten evenly spaced numbers between 1-6. Print all the results below and store them separately:

- i) Cube (i.e. raise to the power of 3) all the elements in both arrays (element-wise)
- ii) Add both the cubed arrays (e.g., $[1,2] + [3,4] = [4,6]$)
- iii) Sum the elements with even indices of the added array.
- iv) Take the square root of the added array (element-wise square root)

```
In [284]: # your code here
import numpy as np
import pandas as pd

a = np.array(range(25,35))
b = np.arange(1,6,.5) #just to make the arrays the same size
print(a)
print(b)

a1 = a ** 3
b1 = b ** 3

print(a1)
print(b1)

c = a1 + b1 #size 11 array and size 10 array but it specifically says in
clusive
print(c)

d = c[::2].sum()
print(d)

f = np.sqrt(c)
print(f)
```

```
[25 26 27 28 29 30 31 32 33 34]
[1.  1.5 2.  2.5 3.  3.5 4.  4.5 5.  5.5]
[15625 17576 19683 21952 24389 27000 29791 32768 35937 39304]
[ 1.      3.375  8.      15.625 27.      42.875 64.      91.125 125.
 166.375]
[15626.      17579.375 19691.      21967.625 24416.      27042.875 29855.
 32859.125 36062.      39470.375]
125650.0
[125.00399994 132.58723543 140.32462364 148.21479346 156.25619988
 164.44717997 172.7859948 181.27086087 189.89997367 198.67152539]
```

1b) Append b to a, reshape the appended array so that it is a 4x5, 2d array and store the results in a variable called m. Print m.

```
In [285]: # your code here
m = np.vstack([a[:5], a[5:], b[:5], b[5:]])
m.T
```

```
Out[285]: array([[25. , 30. , 1. , 3.5],
 [26. , 31. , 1.5, 4. ],
 [27. , 32. , 2. , 4.5],
 [28. , 33. , 2.5, 5. ],
 [29. , 34. , 3. , 5.5]])
```

1c) Extract the third and the fourth column of the m matrix. Store the resulting 4x2 matrix in a new variable called m2. Print m2.

```
In [286]: # your code here
m2 = m[:, [3,4]]
m2
```

```
Out[286]: array([[28. , 29. ],
                [33. , 34. ],
                [ 2.5,  3. ],
                [ 5. ,  5.5]])
```

1d) Take the dot product of m2 and m store the results in a matrix called m3. Print m3. Note that Dot product of two matrices $A.B = A^T B$

```
In [287]: # your code here
m3 = a.T * b
m3
```

```
Out[287]: array([ 25.,  39.,  54.,  70.,  87., 105., 124., 144., 165., 187.])
```

Numpy conditions

2a) Create a numpy array called 'f' where the values are cosine(x) for x from 0 to pi with 50 equally spaced values in f

- Print f
- Use condition on the array and print an array that is True when $f \geq 1/2$ and False when $f < 1/2$
- Create and print an array sequence that has only those values where $f \geq 1/2$

```
In [288]: # your code here
temp = np.arange(0.0, np.pi, np.pi/50, dtype=float)
f = np.cos(temp)
np.where(f >= .5, True, False)
```

```
Out[288]: array([ True,  True,  True,  True,  True,  True,  True,  True,  True,  True,
                True,  True,  True,  True,  True,  True,  True,  True,  True,  True,
                False, False, False, False, False, False, False, False, False, False,
                False, False, False, False, False, False, False, False, False, False,
                False, False, False, False, False, False, False, False, False, False,
                False, False, False, False, False])
```

NumPy and 2 Variable Prediction

Let 'x' be the number of miles a person drives per day and 'y' be the dollars spent on buying car fuel (per day).

We have created 2 numpy arrays each of size 100 that represent x and y.

x (number of miles) ranges from 1 to 10 with a uniform noise of (0,1/2)

y (money spent in dollars) will be from 1 to 20 with a uniform noise (0,1)

```
In [289]: # seed the random number generator with a fixed value
import numpy as np
np.random.seed(500)

x = np.linspace(1,10,100)+ np.random.uniform(low=0,high=.5,size=100)
y = np.linspace(1,20,100)+ np.random.uniform(low=0,high=1,size=100)
print('x = ',x)
print('y = ',y)
```

```

x = [ 1.34683976  1.12176759  1.51512398  1.55233174  1.40619168  1.65
075498
  1.79399331  1.80243817  1.89844195  2.00100023  2.3344038  2.2242487
2
  2.24914511  2.36268477  2.49808849  2.8212704  2.68452475  2.6822942
7
  3.09511169  2.95703884  3.09047742  3.2544361  3.41541904  3.4088637
5
  3.50672677  3.74960644  3.64861355  3.7721462  3.56368566  4.0109270
1
  4.15630694  4.06088549  4.02517179  4.25169402  4.15897504  4.2683533
3
  4.32520644  4.48563164  4.78490721  4.84614839  4.96698768  5.1875425
9
  5.29582013  5.32097781  5.0674106  5.47601124  5.46852704  5.6453745
2
  5.49642807  5.89755027  5.68548923  5.76276141  5.94613234  6.1813571
3
  5.96522091  6.0275473  6.54290191  6.4991329  6.74003765  6.8180980
7
  6.50611821  6.91538752  7.01250925  6.89905417  7.31314433  7.2047229
7
  7.1043621  7.48199528  7.58957227  7.61744354  7.6991707  7.8543682
2
  8.03510784  7.80787781  8.22410224  7.99366248  8.40581097  8.2891379
2
  8.45971515  8.54227144  8.6906456  8.61856507  8.83489887  8.6630965
8
  8.94837987  9.20890222  8.9614749  8.92608294  9.13231416  9.5588989
6
  9.61488451  9.54252979  9.42015491  9.90952569  10.00659591  10.0250426
5
  10.07330937  9.93489915  10.0892334  10.36509991]
y = [ 1.6635012  2.0214592  2.10816052  2.26016496  1.96287558  2.95
54635
  3.02881887  3.33565296  2.75465779  3.4250107  3.39670148  3.3937776
7
  3.78503343  4.38293049  4.32963586  4.03925039  4.73691868  4.3009839
9
  4.8416329  4.78175957  4.99765787  5.31746817  5.76844671  5.9372374
9
  5.72811642  6.70973615  6.68143367  6.57482731  7.17737603  7.5486325
2
  7.30221419  7.3202573  7.78023884  7.91133365  8.2765417  8.6920328
1
  8.78219865  8.45897546  8.89094715  8.81719921  8.87106971  9.6619256
2
  9.4020625  9.85990783  9.60359778  10.07386266  10.6957995  10.6672191
6
  11.18256285  10.57431836  11.46744716  10.94398916  11.26445259  12.0975482
8
  12.11988037  12.121557  12.17613693  12.43750193  13.00912372  12.8640719
4
  13.24640866  12.76120085  13.11723062  14.07841099  14.19821707  14.2728900
1
  14.30624942  14.63060835  14.2770918  15.0744923  14.45261619  15.1189731
3

```

```

15.2378667 15.27203124 15.32491892 16.01095271 15.71250558 16.2948850
6
16.70618934 16.56555394 16.42379457 17.18144744 17.13813976 17.6961362
5
17.37763019 17.90942839 17.90343733 18.01951169 18.35727914 18.1684126
9
18.61813748 18.66062754 18.81217983 19.44995194 19.7213867 19.7196672
6
19.78961904 19.64385088 20.69719809 20.07974319]

```

3a) Find Expected value of x and the expected value of y

```

In [290]: # your code here
print(np.mean(x))
print(np.mean(y))

```

```

5.782532541587923
11.012981683344968

```

3b) Find variance of distributions of x and y

```

In [291]: # your code here
print(np.var(x))
print(np.var(y))

```

```

7.03332752947585
30.113903575509635

```

3c) Find co-variance of x and y.

```

In [292]: # your code here
print(np.cov(x,y))
print(np.cov(x,y)[1][0])

```

```

[[ 7.10437124 14.65774383]
 [14.65774383 30.41808442]]
14.657743832803437

```

3d) Assuming that number of dollars spent in car fuel is only dependant on the miles driven, by a linear relationship.

Write code that uses a linear predictor to calculate a predicted value of y for each x i.e $y_{\text{predicted}} = f(x) = y_0 + mx$. (Do not use Machine learning libraries)

```

In [293]: # your code here
m,b = np.polyfit(x,y,1)
print(m,b)
y_predicted = m*x + b

```

```

2.063200715971324 -0.9175435965867221

```

3e) Predict y for each value in x, put the error into an array called y_error

```
In [294]: # your code here
y_error = [y_predicted for b in x ][0]
y_error
```

```
Out[294]: array([ 1.86125717,  1.39688809,  2.20846128,  2.28522836,  1.98371207,
 2.48829527,  2.78382468,  2.80124813,  2.9993232 ,  3.21092152,
 3.8988      ,  3.67152796,  3.7228942 ,  3.9571493 ,  4.23651436,
 4.9033035 ,  4.62116978,  4.61656787,  5.46829307,  5.18342105,
 5.45873164,  5.79701128,  6.12915141,  6.11562653,  6.31753758,
 6.81864709,  6.61027849,  6.86515115,  6.43505522,  7.35780389,
 7.65775187,  7.46087825,  7.38719373,  7.85455455,  7.66325667,
 7.88892606,  8.00622544,  8.33721481,  8.95468038,  9.08103323,
 9.33034895,  9.78539799, 10.00879629, 10.06070164,  9.53754157,
10.38056671, 10.36512531, 10.72999716, 10.42269073, 11.25028634,
10.81276185, 10.97218988, 11.35052091, 11.83583685, 11.38990445,
11.51849632, 12.58177632, 12.49147206, 12.98850691, 13.14956122,
12.50588416, 13.35028889, 13.5506705 , 13.31658991, 14.17094102,
13.947246  , 13.74018137, 14.51931443, 14.74126735, 14.79877137,
14.96739089, 15.28759454, 15.66049665, 15.1916755 , 16.05043004,
15.57498655, 16.42533161, 16.18461169, 16.53654675, 16.70687695,
17.01300263, 16.86428603, 17.31062607, 16.95616347, 17.54476017,
18.08227006, 17.57177784, 17.49875711, 17.92425351, 18.80438359,
18.91989301, 18.77061069, 18.51812677, 19.5277969 , 19.72807224,
19.76613158, 19.8657155 , 19.58014745, 19.89856998, 20.4677379
7])
```

3f) Write code that calculates the root mean square error(RMSE), that is root of average of y-error squared__

```
In [295]: # your code here
np.sqrt(np.mean((y_error-y)**2))
```

```
Out[295]: 0.4176777236685612
```

Pandas Introduction

Reading File

```
In [296]: # Load required modules
import pandas as pd
import numpy as np
```


Read the CSV file called 'data3.csv' into a dataframe called df.

Data description

- File location: https://bcourses.berkeley.edu/files/74463396/download?download_frd=1
(https://bcourses.berkeley.edu/files/74463396/download?download_frd=1)
- Data source: http://www.fao.org/nr/water/aquastat/data/query/index.html?*lang=en
(http://www.fao.org/nr/water/aquastat/data/query/index.html?*lang=en)
- Data, units:
- GDP, current USD (CPI adjusted)
- NRI, mm/yr
- Population density, inhab/km²
- Total area of the country, 1000 ha = 10km²
- Total Population, unit 1000 inhabitants

```
In [297]: # your code here
file = "data3.csv"
data = pd.read_csv(file)
```

4a) Display the first 10 rows of the dataframe

```
In [298]: # your code here
data.head(10)
```

Out[298]:

| | Area | Area Id | Variable Name | Variable Id | Year | Value | Symbol | Other |
|---|-----------|---------|---------------------------|-------------|--------|----------|--------|-------|
| 0 | Argentina | 9.0 | Total area of the country | 4100.0 | 1962.0 | 278040.0 | E | NaN |
| 1 | Argentina | 9.0 | Total area of the country | 4100.0 | 1967.0 | 278040.0 | E | NaN |
| 2 | Argentina | 9.0 | Total area of the country | 4100.0 | 1972.0 | 278040.0 | E | NaN |
| 3 | Argentina | 9.0 | Total area of the country | 4100.0 | 1977.0 | 278040.0 | E | NaN |
| 4 | Argentina | 9.0 | Total area of the country | 4100.0 | 1982.0 | 278040.0 | E | NaN |
| 5 | Argentina | 9.0 | Total area of the country | 4100.0 | 1987.0 | 278040.0 | E | NaN |
| 6 | Argentina | 9.0 | Total area of the country | 4100.0 | 1992.0 | 278040.0 | E | NaN |
| 7 | Argentina | 9.0 | Total area of the country | 4100.0 | 1997.0 | 278040.0 | E | NaN |
| 8 | Argentina | 9.0 | Total area of the country | 4100.0 | 2002.0 | 278040.0 | E | NaN |
| 9 | Argentina | 9.0 | Total area of the country | 4100.0 | 2007.0 | 278040.0 | E | NaN |

4b) Display the column names.

```
In [299]: # your code here
data.columns
```

```
Out[299]: Index(['Area', 'Area Id', 'Variable Name', 'Variable Id', 'Year', 'Value',
               'Symbol', 'Other'],
              dtype='object')
```

4c) Use `iloc` to display the first 3 rows and first 4 columns.

```
In [300]: # your code here
data.iloc[0:3, 0:4]
```

Out[300]:

| | Area | Area Id | Variable Name | Variable Id |
|---|-----------|---------|---------------------------|-------------|
| 0 | Argentina | 9.0 | Total area of the country | 4100.0 |
| 1 | Argentina | 9.0 | Total area of the country | 4100.0 |
| 2 | Argentina | 9.0 | Total area of the country | 4100.0 |

Data Preprocessing

5a) Find all the rows that have 'NaN' in the 'Symbol' column. Display first 5 rows.

Hint : You might have to use a condition (mask)

```
In [301]: # your code here
#null_col = data.columns[data.isnull().any()]
print(data[data['Symbol'].isnull()][null_columns].head(5))
```

| | Area | Area Id | Variable Name | Variable Id | Year | V |
|-----|-----------------------|---------|---------------|-------------|------|---|
| 390 | NaN | NaN | NaN | NaN | NaN | |
| 391 | E - External data | NaN | NaN | NaN | NaN | |
| 392 | I - AQUASTAT estimate | NaN | NaN | NaN | NaN | |
| 393 | K - Aggregate data | NaN | NaN | NaN | NaN | |
| 394 | L - Modelled data | NaN | NaN | NaN | NaN | |

| | Symbol | Other |
|-----|--------|-------|
| 390 | NaN | NaN |
| 391 | NaN | NaN |
| 392 | NaN | NaN |
| 393 | NaN | NaN |
| 394 | NaN | NaN |

5b) Now, we will try to get rid of the NaN valued rows and columns. First, drop the column 'Other' which only has 'NaN' values. Then drop all other rows that have any column with a value 'NaN'. Then display the last 5 rows of the dataframe.

```
In [302]: # your code here
data = data.drop(columns = 'Other')
```

```
In [303]: null_row = data[data.isnull().any(axis=1)]
null_row.index
data = data.drop(null_row.index)
data.tail(5)
```

Out[303]:

| | Area | Area Id | Variable Name | Variable Id | Year | Value | Symbol |
|-----|--------------------------|---------|-------------------------------|-------------|--------|--------|--------|
| 385 | United States of America | 231.0 | National Rainfall Index (NRI) | 4472.0 | 1981.0 | 949.2 | E |
| 386 | United States of America | 231.0 | National Rainfall Index (NRI) | 4472.0 | 1984.0 | 974.6 | E |
| 387 | United States of America | 231.0 | National Rainfall Index (NRI) | 4472.0 | 1992.0 | 1020.0 | E |
| 388 | United States of America | 231.0 | National Rainfall Index (NRI) | 4472.0 | 1996.0 | 1005.0 | E |
| 389 | United States of America | 231.0 | National Rainfall Index (NRI) | 4472.0 | 2002.0 | 938.7 | E |

6a) For our analysis we do not want all the columns in our dataframe. Lets drop all the redundant columns/ features.

Drop columns: Area Id, Variable Id, Symbol. Save the new dataframe as df1. Display the first 5 rows of the new dataframe.

```
In [304]: # your code here
df1 = data.drop(columns = ['Area Id', 'Variable Id', 'Symbol'])
df1.head(5)
```

Out[304]:

| | Area | Variable Name | Year | Value |
|---|-----------|---------------------------|--------|----------|
| 0 | Argentina | Total area of the country | 1962.0 | 278040.0 |
| 1 | Argentina | Total area of the country | 1967.0 | 278040.0 |
| 2 | Argentina | Total area of the country | 1972.0 | 278040.0 |
| 3 | Argentina | Total area of the country | 1977.0 | 278040.0 |
| 4 | Argentina | Total area of the country | 1982.0 | 278040.0 |

6b) Display all the unique values in your new dataframe for columns: Area, Variable Name, Year.

```
In [305]: # your code here
print(df1["Area"].unique())
print(df1["Variable Name"].unique())
print(df1["Year"].unique())

['Argentina' 'Australia' 'Germany' 'Iceland' 'Ireland' 'Sweden'
 'United States of America']
['Total area of the country' 'Total population' 'Population density'
 'Gross Domestic Product (GDP)' 'National Rainfall Index (NRI)']
[1962. 1967. 1972. 1977. 1982. 1987. 1992. 1997. 2002. 2007. 2012. 2014.
 2015. 1963. 1970. 1974. 1978. 1984. 1990. 1964. 1981. 1985. 1996. 2001.
 1969. 1973. 1979. 1993. 1971. 1975. 1986. 1991. 1998. 2000. 1965. 1983.
 1988. 1995.]
```

6c) Convert the year column to pandas datetime. Convert the 'Year' column float values to pandas datetime objects, where each year is represented as the first day of that year. Also display the column and datatype for 'Year' after conversion. For eg: 1962.0 will be represented as 1962-01-01¶

```
In [306]: df1['month'] = 1
df1['day'] = 1
df1['Year'] = pd.to_datetime(df1.drop(columns = ['Area', 'Value', 'Variable Name']))
df1 = df1.drop(columns = ['month', 'day'])
df1
```

Out[306]:

| | Area | Variable Name | Year | Value |
|-----|--------------------------|---------------------------|------------|--------------|
| 0 | Argentina | Total area of the country | 1962-01-01 | 2.780400e+05 |
| 1 | Argentina | Total area of the country | 1967-01-01 | 2.780400e+05 |
| 2 | Argentina | Total area of the country | 1972-01-01 | 2.780400e+05 |
| 3 | Argentina | Total area of the country | 1977-01-01 | 2.780400e+05 |
| 4 | Argentina | Total area of the country | 1982-01-01 | 2.780400e+05 |
| 5 | Argentina | Total area of the country | 1987-01-01 | 2.780400e+05 |
| 6 | Argentina | Total area of the country | 1992-01-01 | 2.780400e+05 |
| 7 | Argentina | Total area of the country | 1997-01-01 | 2.780400e+05 |
| 8 | Argentina | Total area of the country | 2002-01-01 | 2.780400e+05 |
| 9 | Argentina | Total area of the country | 2007-01-01 | 2.780400e+05 |
| 10 | Argentina | Total area of the country | 2012-01-01 | 2.780400e+05 |
| 11 | Argentina | Total area of the country | 2014-01-01 | 2.780400e+05 |
| 12 | Argentina | Total population | 1962-01-01 | 2.128800e+04 |
| 13 | Argentina | Total population | 1967-01-01 | 2.293200e+04 |
| 14 | Argentina | Total population | 1972-01-01 | 2.478300e+04 |
| 15 | Argentina | Total population | 1977-01-01 | 2.687900e+04 |
| 16 | Argentina | Total population | 1982-01-01 | 2.899400e+04 |
| 17 | Argentina | Total population | 1987-01-01 | 3.132600e+04 |
| 18 | Argentina | Total population | 1992-01-01 | 3.365500e+04 |
| 19 | Argentina | Total population | 1997-01-01 | 3.583400e+04 |
| 20 | Argentina | Total population | 2002-01-01 | 3.788900e+04 |
| 21 | Argentina | Total population | 2007-01-01 | 3.997000e+04 |
| 22 | Argentina | Total population | 2012-01-01 | 4.209500e+04 |
| 23 | Argentina | Total population | 2015-01-01 | 4.341700e+04 |
| 24 | Argentina | Population density | 1962-01-01 | 7.656000e+00 |
| 25 | Argentina | Population density | 1967-01-01 | 8.248000e+00 |
| 26 | Argentina | Population density | 1972-01-01 | 8.913000e+00 |
| 27 | Argentina | Population density | 1977-01-01 | 9.667000e+00 |
| 28 | Argentina | Population density | 1982-01-01 | 1.043000e+01 |
| 29 | Argentina | Population density | 1987-01-01 | 1.127000e+01 |
| ... | ... | ... | ... | ... |
| 360 | United States of America | Population density | 1972-01-01 | 2.214000e+01 |
| 361 | United States of America | Population density | 1977-01-01 | 2.317000e+01 |
| 362 | United States of America | Population density | 1982-01-01 | 2.430000e+01 |

| | Area | Variable Name | Year | Value |
|-----|--------------------------|-------------------------------|------------|--------------|
| 363 | United States of America | Population density | 1987-01-01 | 2.549000e+01 |
| 364 | United States of America | Population density | 1992-01-01 | 2.678000e+01 |
| 365 | United States of America | Population density | 1997-01-01 | 2.834000e+01 |
| 366 | United States of America | Population density | 2002-01-01 | 2.995000e+01 |
| 367 | United States of America | Population density | 2007-01-01 | 3.132000e+01 |
| 368 | United States of America | Population density | 2012-01-01 | 3.202000e+01 |
| 369 | United States of America | Population density | 2015-01-01 | 3.273000e+01 |
| 370 | United States of America | Gross Domestic Product (GDP) | 1962-01-01 | 6.050000e+11 |
| 371 | United States of America | Gross Domestic Product (GDP) | 1967-01-01 | 8.620000e+11 |
| 372 | United States of America | Gross Domestic Product (GDP) | 1972-01-01 | 1.280000e+12 |
| 373 | United States of America | Gross Domestic Product (GDP) | 1977-01-01 | 2.090000e+12 |
| 374 | United States of America | Gross Domestic Product (GDP) | 1982-01-01 | 3.340000e+12 |
| 375 | United States of America | Gross Domestic Product (GDP) | 1987-01-01 | 4.870000e+12 |
| 376 | United States of America | Gross Domestic Product (GDP) | 1992-01-01 | 6.540000e+12 |
| 377 | United States of America | Gross Domestic Product (GDP) | 1997-01-01 | 8.610000e+12 |
| 378 | United States of America | Gross Domestic Product (GDP) | 2002-01-01 | 1.100000e+13 |
| 379 | United States of America | Gross Domestic Product (GDP) | 2007-01-01 | 1.450000e+13 |
| 380 | United States of America | Gross Domestic Product (GDP) | 2012-01-01 | 1.620000e+13 |
| 381 | United States of America | Gross Domestic Product (GDP) | 2015-01-01 | 1.790000e+13 |
| 382 | United States of America | National Rainfall Index (NRI) | 1965-01-01 | 9.285000e+02 |
| 383 | United States of America | National Rainfall Index (NRI) | 1969-01-01 | 9.522000e+02 |
| 384 | United States of America | National Rainfall Index (NRI) | 1974-01-01 | 1.008000e+03 |
| 385 | United States of America | National Rainfall Index (NRI) | 1981-01-01 | 9.492000e+02 |
| 386 | United States of America | National Rainfall Index (NRI) | 1984-01-01 | 9.746000e+02 |
| 387 | United States of America | National Rainfall Index (NRI) | 1992-01-01 | 1.020000e+03 |
| 388 | United States of America | National Rainfall Index (NRI) | 1996-01-01 | 1.005000e+03 |
| 389 | United States of America | National Rainfall Index (NRI) | 2002-01-01 | 9.387000e+02 |

390 rows × 4 columns

Extract specific statistics from the preprocessed data:

7a) Create a dataframe 'dftemp' to store rows where Area is 'Iceland'. Display the dataframe.

```
In [319]: # your code here
          dftemp = (df1[df1['Area'] == "Iceland"])
          dftemp
```


Out[319]:

| | Area | Variable Name | Year | Value |
|-----|---------|---------------------------|------------|--------------|
| 166 | Iceland | Total area of the country | 1962-01-01 | 1.030000e+04 |
| 167 | Iceland | Total area of the country | 1967-01-01 | 1.030000e+04 |
| 168 | Iceland | Total area of the country | 1972-01-01 | 1.030000e+04 |
| 169 | Iceland | Total area of the country | 1977-01-01 | 1.030000e+04 |
| 170 | Iceland | Total area of the country | 1982-01-01 | 1.030000e+04 |
| 171 | Iceland | Total area of the country | 1987-01-01 | 1.030000e+04 |
| 172 | Iceland | Total area of the country | 1992-01-01 | 1.030000e+04 |
| 173 | Iceland | Total area of the country | 1997-01-01 | 1.030000e+04 |
| 174 | Iceland | Total area of the country | 2002-01-01 | 1.030000e+04 |
| 175 | Iceland | Total area of the country | 2007-01-01 | 1.030000e+04 |
| 176 | Iceland | Total area of the country | 2012-01-01 | 1.030000e+04 |
| 177 | Iceland | Total area of the country | 2014-01-01 | 1.030000e+04 |
| 178 | Iceland | Total population | 1962-01-01 | 1.826000e+02 |
| 179 | Iceland | Total population | 1967-01-01 | 1.974000e+02 |
| 180 | Iceland | Total population | 1972-01-01 | 2.099000e+02 |
| 181 | Iceland | Total population | 1977-01-01 | 2.221000e+02 |
| 182 | Iceland | Total population | 1982-01-01 | 2.331000e+02 |
| 183 | Iceland | Total population | 1987-01-01 | 2.469000e+02 |
| 184 | Iceland | Total population | 1992-01-01 | 2.599000e+02 |
| 185 | Iceland | Total population | 1997-01-01 | 2.728000e+02 |
| 186 | Iceland | Total population | 2002-01-01 | 2.869000e+02 |
| 187 | Iceland | Total population | 2007-01-01 | 3.054000e+02 |
| 188 | Iceland | Total population | 2012-01-01 | 3.234000e+02 |
| 189 | Iceland | Total population | 2015-01-01 | 3.294000e+02 |
| 190 | Iceland | Population density | 1962-01-01 | 1.773000e+00 |
| 191 | Iceland | Population density | 1967-01-01 | 1.917000e+00 |
| 192 | Iceland | Population density | 1972-01-01 | 2.038000e+00 |
| 193 | Iceland | Population density | 1977-01-01 | 2.156000e+00 |
| 194 | Iceland | Population density | 1982-01-01 | 2.263000e+00 |
| 195 | Iceland | Population density | 1987-01-01 | 2.397000e+00 |
| 196 | Iceland | Population density | 1992-01-01 | 2.523000e+00 |
| 197 | Iceland | Population density | 1997-01-01 | 2.649000e+00 |
| 198 | Iceland | Population density | 2002-01-01 | 2.785000e+00 |
| 199 | Iceland | Population density | 2007-01-01 | 2.965000e+00 |

| | Area | Variable Name | Year | Value |
|-----|---------|-------------------------------|------------|--------------|
| 200 | Iceland | Population density | 2012-01-01 | 3.140000e+00 |
| 201 | Iceland | Population density | 2015-01-01 | 3.198000e+00 |
| 202 | Iceland | Gross Domestic Product (GDP) | 1962-01-01 | 2.849165e+08 |
| 203 | Iceland | Gross Domestic Product (GDP) | 1967-01-01 | 6.212260e+08 |
| 204 | Iceland | Gross Domestic Product (GDP) | 1972-01-01 | 8.465069e+08 |
| 205 | Iceland | Gross Domestic Product (GDP) | 1977-01-01 | 2.226539e+09 |
| 206 | Iceland | Gross Domestic Product (GDP) | 1982-01-01 | 3.232804e+09 |
| 207 | Iceland | Gross Domestic Product (GDP) | 1987-01-01 | 5.565384e+09 |
| 208 | Iceland | Gross Domestic Product (GDP) | 1992-01-01 | 7.138788e+09 |
| 209 | Iceland | Gross Domestic Product (GDP) | 1997-01-01 | 7.596126e+09 |
| 210 | Iceland | Gross Domestic Product (GDP) | 2002-01-01 | 9.161798e+09 |
| 211 | Iceland | Gross Domestic Product (GDP) | 2007-01-01 | 2.129384e+10 |
| 212 | Iceland | Gross Domestic Product (GDP) | 2012-01-01 | 1.419452e+10 |
| 213 | Iceland | Gross Domestic Product (GDP) | 2015-01-01 | 1.659849e+10 |
| 214 | Iceland | National Rainfall Index (NRI) | 1967-01-01 | 8.160000e+02 |
| 215 | Iceland | National Rainfall Index (NRI) | 1971-01-01 | 9.632000e+02 |
| 216 | Iceland | National Rainfall Index (NRI) | 1975-01-01 | 1.010000e+03 |
| 217 | Iceland | National Rainfall Index (NRI) | 1981-01-01 | 9.326000e+02 |
| 218 | Iceland | National Rainfall Index (NRI) | 1986-01-01 | 9.685000e+02 |
| 219 | Iceland | National Rainfall Index (NRI) | 1991-01-01 | 1.095000e+03 |
| 220 | Iceland | National Rainfall Index (NRI) | 1997-01-01 | 9.932000e+02 |
| 221 | Iceland | National Rainfall Index (NRI) | 1998-01-01 | 9.234000e+02 |

7b) Print the years when the National Rainfall Index (NRI) was greater than 900 and less than 950 in Iceland. Use the dataframe you created in the previous question 'dftemp'.

```
In [329]: # your code here
temp2 = (dftemp[dftemp['Variable Name'] == "National Rainfall Index (NRI)"])
print((temp2[temp2['Value'].between(900, 950)]))
```

| | Area | Variable Name | Year | Value |
|-----|---------|-------------------------------|------------|-------|
| 217 | Iceland | National Rainfall Index (NRI) | 1981-01-01 | 932.6 |
| 221 | Iceland | National Rainfall Index (NRI) | 1998-01-01 | 923.4 |

US statistics:

8a) Create a new DataFrame called `df_usa` that only contains values where 'Area' is equal to 'United States of America'. Set the indices to be the 'Year' column (Use `.set_index()`). Display the dataframe head.

```
In [335]: # your code here
df_usa = (df1[df1['Area'] == "United States of America"]).set_index('Year')
df_usa.head()
```

Out[335]:

| | Area | Variable Name | Value |
|------------|--------------------------|---------------------------|----------|
| Year | | | |
| 1962-01-01 | United States of America | Total area of the country | 962909.0 |
| 1967-01-01 | United States of America | Total area of the country | 962909.0 |
| 1972-01-01 | United States of America | Total area of the country | 962909.0 |
| 1977-01-01 | United States of America | Total area of the country | 962909.0 |
| 1982-01-01 | United States of America | Total area of the country | 962909.0 |

8b) Pivot the DataFrame so that the unique values in the column 'Variable Name' becomes the columns. The DataFrame values should be the ones in the 'Value' column. Save it in `df_usa`. Display the dataframe head.

```
In [340]: # your code here
df_usa = df_usa.pivot(columns = "Variable Name", values = "Value")
df_usa.head()
```

Out[340]:

| | Variable Name | Gross Domestic Product (GDP) | National Rainfall Index (NRI) | Population density | Total area of the country | Total population |
|------------|---------------|------------------------------|-------------------------------|--------------------|---------------------------|------------------|
| Year | | | | | | |
| 1962-01-01 | | 6.050000e+11 | NaN | 19.93 | 962909.0 | 191861.0 |
| 1965-01-01 | | NaN | 928.5 | NaN | NaN | NaN |
| 1967-01-01 | | 8.620000e+11 | NaN | 21.16 | 962909.0 | 203713.0 |
| 1969-01-01 | | NaN | 952.2 | NaN | NaN | NaN |
| 1972-01-01 | | 1.280000e+12 | NaN | 22.14 | 962909.0 | 213220.0 |

8c) Rename new columns to ['GDP','NRI','PD','Area','Population'] and display the head.

```
In [342]: # your code here
df_usa = df_usa.rename(index=str, columns={"Gross Domestic Product (GDP)": "GDP",
                                           "National Rainfall Index (NRI)": "NRI",
                                           "Population density": "PD",
                                           "Total area of the country": "Area",
                                           "Total population": "Population"})
df_usa.head()
```

Out[342]:

| Variable Name | GDP | NRI | PD | Area | Population |
|---------------------|--------------|-------|-------|----------|------------|
| Year | | | | | |
| 1962-01-01 00:00:00 | 6.050000e+11 | NaN | 19.93 | 962909.0 | 191861.0 |
| 1965-01-01 00:00:00 | NaN | 928.5 | NaN | NaN | NaN |
| 1967-01-01 00:00:00 | 8.620000e+11 | NaN | 21.16 | 962909.0 | 203713.0 |
| 1969-01-01 00:00:00 | NaN | 952.2 | NaN | NaN | NaN |
| 1972-01-01 00:00:00 | 1.280000e+12 | NaN | 22.14 | 962909.0 | 213220.0 |

8d) Replace all 'Nan' values in df_usa with 0. Display the head of the dataframe.

```
In [345]: # your code here
df_usa = df_usa.fillna(0)
df_usa.head()
```

Out[345]:

| Variable Name | GDP | NRI | PD | Area | Population |
|---------------------|--------------|-------|-------|----------|------------|
| Year | | | | | |
| 1962-01-01 00:00:00 | 6.050000e+11 | 0.0 | 19.93 | 962909.0 | 191861.0 |
| 1965-01-01 00:00:00 | 0.000000e+00 | 928.5 | 0.00 | 0.0 | 0.0 |
| 1967-01-01 00:00:00 | 8.620000e+11 | 0.0 | 21.16 | 962909.0 | 203713.0 |
| 1969-01-01 00:00:00 | 0.000000e+00 | 952.2 | 0.00 | 0.0 | 0.0 |
| 1972-01-01 00:00:00 | 1.280000e+12 | 0.0 | 22.14 | 962909.0 | 213220.0 |

Note: Use df_usa

9a) Multiply the 'Area' column for all countries by 10 (so instead of 1000 ha, the unit becomes 100 ha = 1km^2). Display the dataframe head.

```
In [347]: # your code here
df_usa['Area'] = df_usa['Area']*10
df_usa.head()
```

Out[347]:

| Variable Name | GDP | NRI | PD | Area | Population |
|---------------------|--------------|-------|-------|-----------|------------|
| Year | | | | | |
| 1962-01-01 00:00:00 | 6.050000e+11 | 0.0 | 19.93 | 9629090.0 | 191861.0 |
| 1965-01-01 00:00:00 | 0.000000e+00 | 928.5 | 0.00 | 0.0 | 0.0 |
| 1967-01-01 00:00:00 | 8.620000e+11 | 0.0 | 21.16 | 9629090.0 | 203713.0 |
| 1969-01-01 00:00:00 | 0.000000e+00 | 952.2 | 0.00 | 0.0 | 0.0 |
| 1972-01-01 00:00:00 | 1.280000e+12 | 0.0 | 22.14 | 9629090.0 | 213220.0 |

9b) Create a new column in df_usa called 'GDP/capita' and populate it with the calculated GDP per capita. Round the results to two decimal points. Display the dataframe head.

GDP per capita = (GDP / Population)

```
In [352]: # your code here
df_usa['GDP/capita'] = df_usa['GDP']/df_usa['Population']
df_usa['GDP/capita'] = df_usa['GDP/capita'].round(2)
df_usa.head()
```

Out[352]:

| Variable Name | GDP | NRI | PD | Area | Population | GDP/capita |
|---------------------|--------------|-------|-------|-----------|------------|------------|
| Year | | | | | | |
| 1962-01-01 00:00:00 | 6.050000e+11 | 0.0 | 19.93 | 9629090.0 | 191861.0 | 3153324.54 |
| 1965-01-01 00:00:00 | 0.000000e+00 | 928.5 | 0.00 | 0.0 | 0.0 | NaN |
| 1967-01-01 00:00:00 | 8.620000e+11 | 0.0 | 21.16 | 9629090.0 | 203713.0 | 4231443.26 |
| 1969-01-01 00:00:00 | 0.000000e+00 | 952.2 | 0.00 | 0.0 | 0.0 | NaN |
| 1972-01-01 00:00:00 | 1.280000e+12 | 0.0 | 22.14 | 9629090.0 | 213220.0 | 6003189.19 |

9c) Find the maximum value of the 'NRI' column in the US (using pandas methods). What year does the max value occur? Display the values.

```
In [354]: # your code here
print(df_usa['NRI'].max())
df_usa.loc[df_usa['NRI'].idxmax()]
```

1020.0

```
Out[354]: Variable Name
GDP          6.540000e+12
NRI           1.020000e+03
PD            2.678000e+01
Area          9.629090e+06
Population    2.579080e+05
GDP/capita    2.535788e+07
Name: 1992-01-01 00:00:00, dtype: float64
```

```
In [ ]:
```